

CS5525 Final Project Code Submission

Jennifer Appiah-Kubi, Rebecca DeSipio, Ajinkya Fotedar

12/11/2021

Contents

Classification Methods	1
Decision Trees	2
KNN	10
SVM	11
Logistic Regression	16

Classification Methods

First set the directory (path which contains the `heart.csv` data), and import any needed libraries

```
#setwd(" ") # uncomment to set working directory via code
library(tree)
library(randomForest) # bootstrap/bagging & random forest
```

```
## randomForest 4.6-14
```

```
## Type rfNews() to see new features/changes/bug fixes.
```

```
library(class) # KNN
```

```
## Warning: package 'class' was built under R version 4.1.2
```

```
library(caret) # SVM
```

```
## Warning: package 'caret' was built under R version 4.1.2
```

```
## Loading required package: ggplot2
```

```
##
```

```
## Attaching package: 'ggplot2'
```

```
## The following object is masked from 'package:randomForest':
##
##     margin
```

```
## Loading required package: lattice
```

```
library(glmnet)           # logistic regression
```

```
## Loading required package: Matrix
```

```
## Loaded glmnet 4.1-2
```

Decision Trees

```
# Read in and organize data
## -- Read data
heart <- read.csv("heart.csv")
Target <- as.factor(heart$target) # target heart rate


## -- Split into training and test sets
train <- sample(1:nrow(heart), 0.75*nrow(heart))
heart.test <- heart[-train, ]
Target.test <- Target[-train]

# ----- #
#                               Fit a Classification Tree                               #
# ----- #

# Fit a classification tree to the training data
set.seed(2441139)
tree.heart <- tree(Target~. -target, heart, subset=train)
summary(tree.heart)

##
## Classification tree:
## tree(formula = Target ~ . - target, data = heart, subset = train)
## Variables actually used in tree construction:
##  [1] "cp"      "ca"      "thal"    "i.age"   "chol"    "oldpeak"
##  [7] "exang"   "trestbps" "sex"     "slope"   "thalach"
## Number of terminal nodes: 21
## Residual mean deviance: 0.379 = 78.07 / 206
## Misclassification error rate: 0.1013 = 23 / 227

## -- Plot tree
plot(tree.heart)
text(tree.heart, pretty=1, cex=0.7)
```



```
# Prune the classification tree
```

```
set.seed(2441139)
```

```
cv.heart <- cv.tree(tree.heart, FUN=prune.misclass)
```

```
cv.heart
```

```
## $size
```

```
## [1] 21 12 10 4 2 1
```

```
##
```

```
## $dev
```

```
## [1] 52 48 44 39 52 102
```

```
##
```

```
## $k
```

```
## [1] -Inf 0.000000 1.000000 1.833333 7.500000 51.000000
```

```
##
```

```
## $method
```


```
## [1] "misclass"
```

```
##
```


```
## attr(,"class")
```

```
## [1] "prune" "tree.sequence"
```

```
plot(cv.heart$size, cv.heart$dev, type='b')
```



```
prune.heart <- prune.misclass(tree.heart, best=10)  
plot(prune.heart)  
text(prune.heart, pretty=1, cex=0.65)
```



```
# Predict using test set and pruned tree. Compare.
```

```
tree.pred <- predict(tree.heart, heart.test, type='class') # test tree
prune.pred <- predict(prune.heart, heart.test, type='class') # pruned tree
```

```
table(prune.pred, Target.test)
```

```
##           Target.test
## prune.pred 0  1
##           0 24  6
##           1 12 34
```

```
table(tree.pred, Target.test)
```

```
##           Target.test
## tree.pred 0  1
##           0 25 10
##           1 11 30
```

```
# ----- #
#                               Bagging                               #
# ----- #
set.seed(2441139)
```

```
# Perform bagging
```

```
bag.heart <- randomForest(as.factor(as.character(heart$target))~., data=heart,
                          subset=train, mtry=ncol(heart)-1,
                          importance=TRUE)

bag.heart
```

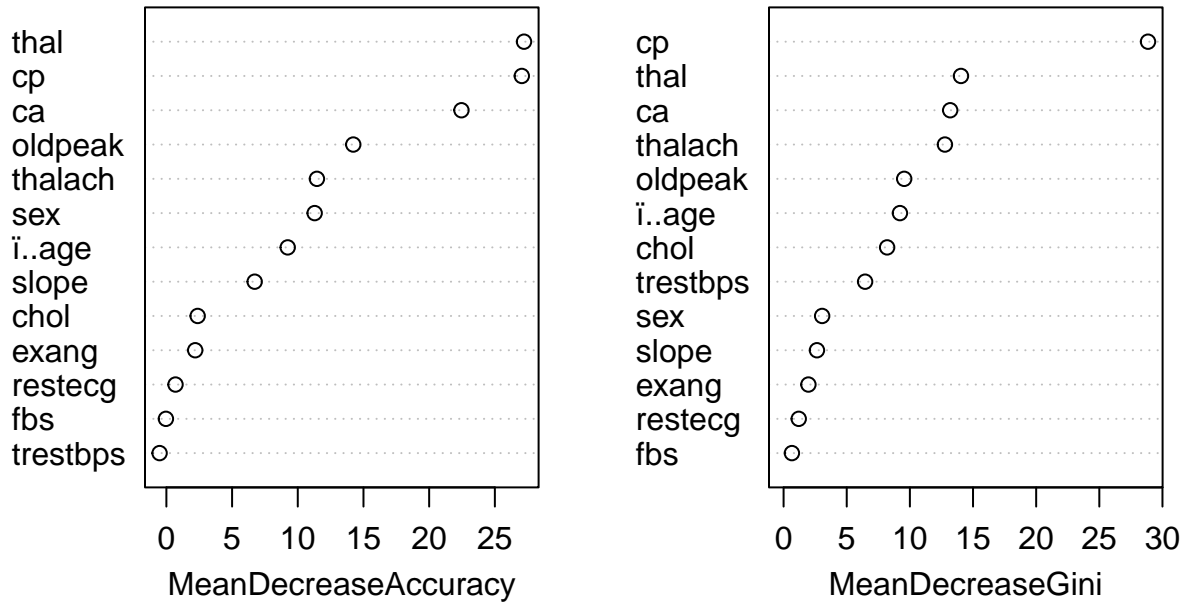
```
##
## Call:
## randomForest(formula = as.factor(as.character(heart$target)) ~ ., data = heart, mtry = ncol(heart),
##              Type of random forest: classification
##              Number of trees: 500
## No. of variables tried at each split: 13
##
## OOB estimate of error rate: 19.82%
## Confusion matrix:
##      0   1 class.error
## 0 77  25    0.245098
## 1 20 105    0.160000
```

```
# Predict on bagged tree
bag.pred <- predict(bag.heart, heart.test, type='class')
table(bag.pred, Target.test)
```

```
##           Target.test
## bag.pred  0   1
##           0 25  6
##           1 11 34
```

```
varImpPlot(bag.heart)
```

bag.heart



```
# -----#
# Random Forest#
# -----#
set.seed(2441139)

# Perform Random Forest
rf.heart <- randomForest(as.factor(as.character(heart$target))~., data=heart,
                          subset=train, mtry=sqrt(ncol(heart)-1),
                          ntree=25, importance=TRUE)

rf.heart
```

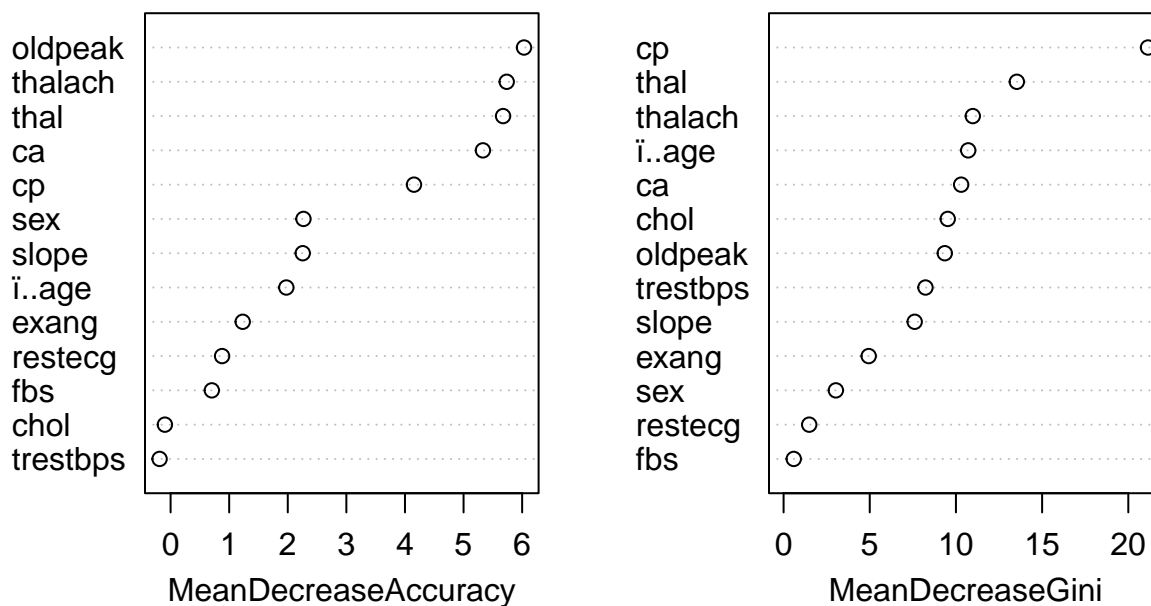
```
##
## Call:
## randomForest(formula = as.factor(as.character(heart$target)) ~ ., data = heart, mtry = sqrt(ncol(heart)-1),
##               Type of random forest: classification
##               Number of trees: 25
##               No. of variables tried at each split: 4
##
##               OOB estimate of error rate: 20.26%
## Confusion matrix:
##      0   1 class.error
## 0 76  26   0.254902
## 1 20 105   0.160000
```

```
# Predict on the forest
rf.pred <- predict(rf.heart, heart.test, type='class')
table(rf.pred, Target.test)
```

```
##      Target.test
## rf.pred  0  1
##      0 30  7
##      1  6 33
```

```
varImpPlot(rf.heart)
```

rf.heart




```
# ----- #
# ----- #
# Determine Best Model (Random Forest) #
# ----- #
# ----- #
# Investigate how mtry affect the accuracy
Acc <- rep(0,ncol(heart)-2)
for (m in 1:(ncol(heart)-2)){
  set.seed(2441139)
  rf.heart <- randomForest(as.factor(as.character(heart$target))~., data=heart,
                           subset=train, mtry=m,
                           ntree=25)
  rf.pred <- predict(rf.heart, heart.test, type='class')
```



```

t <- table(rf.pred, Target.test)
acc <- sum(diag(t))/sum(t)
Acc[m] <- acc
}
mbest <- which(Acc==max(Acc))
plot(1:(ncol(heart)-2), Acc, xlab='mtry', ylab='Accuracy of random forest') # include plot in final sub

```



```

# Now use the best value of m for the random forest
set.seed(2441139)
rf.heart <- randomForest(as.factor(as.character(heart$target))~., data=heart,
                          subset=train, mtry=mbest,
                          ntree=25, importance=TRUE)
rf.heart

##
## Call:
## randomForest(formula = as.factor(as.character(heart$target)) ~ ., data = heart, mtry = mbest,
##               Type of random forest: classification
##               Number of trees: 25
## No. of variables tried at each split: 1
##
## OOB estimate of error rate: 22.03%
## Confusion matrix:
##      0      1 class.error
## 0 76  26      0.254902

```

```
## 1 24 101    0.192000
```



```
# Predict on the forest
```

```
rf.pred <- predict(rf.heart, heart.test, type='class')
table(rf.pred, Target.test)
```

```
##          Target.test
## rf.pred  0  1
##          0 29  5
##          1  7 35
```

```
varImpPlot(rf.heart)
```

rf.heart



KNN

```
# K-Nearest Neighbor
```

```
cl <- as.factor(heart$target[train])
knn.heart <- knn(heart[train,], heart.test, cl, k = 5, prob=TRUE)
table(knn.heart, Target.test)
```

```
##          Target.test
## knn.heart  0  1
##          0 18  7
##          1 18 33
```

SVM

```
set.seed(2441139)

# splitting data into test and train
intrain <- createDataPartition(y = heart$target, p = 0.7, list = F)

training <- heart[intrain,]
testing <- heart[-intrain,]
training[["target"]] <- as.factor(training[["target"]])

dim(training)

## [1] 213 14

dim(testing)

## [1] 90 14

# model training with svm
trctrl <- trainControl(method = "repeatedcv", number = 10, repeats = 3)
svm.mod <- train(target ~ ., data = training, method = "svmLinear",
                 trControl = trctrl,
                 preProcess = c("center", "scale"),
                 tuneLength = 10)

svm.mod

## Support Vector Machines with Linear Kernel
##
## 213 samples
## 13 predictor
## 2 classes: '0', '1'
##
## Pre-processing: centered (13), scaled (13)
## Resampling: Cross-Validated (10 fold, repeated 3 times)
## Summary of sample sizes: 192, 191, 192, 192, 191, 191, ...
## Resampling results:
##
## Accuracy Kappa
## 0.8266955 0.6511675
##
## Tuning parameter 'C' was held constant at a value of 1

# prediction using the above model
svm.pred <- predict(svm.mod, newdata = testing)
svm.pred

## [1] 1 1 1 1 1 1 1 0 1 1 1 1 1 1 1 1 0 0 0 1 0 1 1 1 1 1 1 0 1 1 0 0 1 1 1 1 1
## [39] 1 1 1 1 1 1 1 0 0 1 1 1 1 0 1 0 1 1 1 0 0 0 0 0 0 0 0 0 0 0 1 0 0 0 0 0 0 0
## [77] 0 0 0 1 1 1 0 0 1 0 1 0 0 0
## Levels: 0 1
```

```
# accuracy of the trained model
confusionMatrix(table(svm.pred, testing$target))
```

```
## Confusion Matrix and Statistics
##
##
## svm.pred  0  1
##          0 28 11
##          1  9 42
##
##              Accuracy : 0.7778
##              95% CI : (0.6779, 0.8587)
##      No Information Rate : 0.5889
##      P-Value [Acc > NIR] : 0.0001266
##
##              Kappa : 0.5448
##
##  Mcnemar's Test P-Value : 0.8230633
##
##      Sensitivity : 0.7568
##      Specificity : 0.7925
##      Pos Pred Value : 0.7179
##      Neg Pred Value : 0.8235
##      Prevalence : 0.4111
##      Detection Rate : 0.3111
##      Detection Prevalence : 0.4333
##      Balanced Accuracy : 0.7746
##
##      'Positive' Class : 0
##
```

```
# costs for further tuning with 10-fold cross-validation
grid <- expand.grid(C = c(0, 0.01, 0.05, 0.1, 0.25, 0.5, 0.75, 1, 1.25, 1.5, 1.75, 2, 5))

svm.mod.grid <- train(target ~ ., data = training, method = "svmLinear",
                      trControl = trctrl,
                      preProcess = c("center", "scale"),
                      tuneGrid = grid,
                      tuneLength = 10)
```

```
## Warning: model fit failed for Fold01.Rep1: C=0.00 Error in .local(x, ...) :
##   No Support Vectors found. You may want to change your parameters
```

```
## Warning: model fit failed for Fold02.Rep1: C=0.00 Error in .local(x, ...) :
##   No Support Vectors found. You may want to change your parameters
```

```
## Warning: model fit failed for Fold03.Rep1: C=0.00 Error in .local(x, ...) :
##   No Support Vectors found. You may want to change your parameters
```

```
## Warning: model fit failed for Fold04.Rep1: C=0.00 Error in .local(x, ...) :
##   No Support Vectors found. You may want to change your parameters
```

```
## Warning: model fit failed for Fold05.Rep1: C=0.00 Error in .local(x, ...) :  
##   No Support Vectors found. You may want to change your parameters  
  
## Warning: model fit failed for Fold06.Rep1: C=0.00 Error in .local(x, ...) :  
##   No Support Vectors found. You may want to change your parameters  
  
## Warning: model fit failed for Fold07.Rep1: C=0.00 Error in .local(x, ...) :  
##   No Support Vectors found. You may want to change your parameters  
  
## Warning: model fit failed for Fold08.Rep1: C=0.00 Error in .local(x, ...) :  
##   No Support Vectors found. You may want to change your parameters  
  
## Warning: model fit failed for Fold09.Rep1: C=0.00 Error in .local(x, ...) :  
##   No Support Vectors found. You may want to change your parameters  
  
## Warning: model fit failed for Fold10.Rep1: C=0.00 Error in .local(x, ...) :  
##   No Support Vectors found. You may want to change your parameters  
  
## Warning: model fit failed for Fold01.Rep2: C=0.00 Error in .local(x, ...) :  
##   No Support Vectors found. You may want to change your parameters  
  
## Warning: model fit failed for Fold02.Rep2: C=0.00 Error in .local(x, ...) :  
##   No Support Vectors found. You may want to change your parameters  
  
## Warning: model fit failed for Fold03.Rep2: C=0.00 Error in .local(x, ...) :  
##   No Support Vectors found. You may want to change your parameters  
  
## Warning: model fit failed for Fold04.Rep2: C=0.00 Error in .local(x, ...) :  
##   No Support Vectors found. You may want to change your parameters  
  
## Warning: model fit failed for Fold05.Rep2: C=0.00 Error in .local(x, ...) :  
##   No Support Vectors found. You may want to change your parameters  
  
## Warning: model fit failed for Fold06.Rep2: C=0.00 Error in .local(x, ...) :  
##   No Support Vectors found. You may want to change your parameters  
  
## Warning: model fit failed for Fold07.Rep2: C=0.00 Error in .local(x, ...) :  
##   No Support Vectors found. You may want to change your parameters  
  
## Warning: model fit failed for Fold08.Rep2: C=0.00 Error in .local(x, ...) :  
##   No Support Vectors found. You may want to change your parameters  
  
## Warning: model fit failed for Fold09.Rep2: C=0.00 Error in .local(x, ...) :  
##   No Support Vectors found. You may want to change your parameters  
  
## Warning: model fit failed for Fold10.Rep2: C=0.00 Error in .local(x, ...) :  
##   No Support Vectors found. You may want to change your parameters  
  
## Warning: model fit failed for Fold01.Rep3: C=0.00 Error in .local(x, ...) :  
##   No Support Vectors found. You may want to change your parameters
```

```

## Warning: model fit failed for Fold02.Rep3: C=0.00 Error in .local(x, ...) :
##   No Support Vectors found. You may want to change your parameters

## Warning: model fit failed for Fold03.Rep3: C=0.00 Error in .local(x, ...) :
##   No Support Vectors found. You may want to change your parameters

## Warning: model fit failed for Fold04.Rep3: C=0.00 Error in .local(x, ...) :
##   No Support Vectors found. You may want to change your parameters

## Warning: model fit failed for Fold05.Rep3: C=0.00 Error in .local(x, ...) :
##   No Support Vectors found. You may want to change your parameters

## Warning: model fit failed for Fold06.Rep3: C=0.00 Error in .local(x, ...) :
##   No Support Vectors found. You may want to change your parameters

## Warning: model fit failed for Fold07.Rep3: C=0.00 Error in .local(x, ...) :
##   No Support Vectors found. You may want to change your parameters

## Warning: model fit failed for Fold08.Rep3: C=0.00 Error in .local(x, ...) :
##   No Support Vectors found. You may want to change your parameters

## Warning: model fit failed for Fold09.Rep3: C=0.00 Error in .local(x, ...) :
##   No Support Vectors found. You may want to change your parameters

## Warning: model fit failed for Fold10.Rep3: C=0.00 Error in .local(x, ...) :
##   No Support Vectors found. You may want to change your parameters

## Warning in nominalTrainWorkflow(x = x, y = y, wts = weights, info = trainInfo, :
## There were missing values in resampled performance measures.

## Warning in train.default(x, y, weights = w, ...): missing values found in
## aggregated results

```

```
svm.mod.grid
```


```

## Support Vector Machines with Linear Kernel
##
## 213 samples
## 13 predictor
## 2 classes: '0', '1'
##
## Pre-processing: centered (13), scaled (13)
## Resampling: Cross-Validated (10 fold, repeated 3 times)
## Summary of sample sizes: 192, 191, 192, 192, 192, 192, ...
## Resampling results across tuning parameters:
##
##   C      Accuracy   Kappa
## 0.00      NaN      NaN
## 0.01 0.8235209 0.6414705
## 0.05 0.8359307 0.6677455
## 0.10 0.8374459 0.6712576

```

```
## 0.25 0.8437229 0.6843572
## 0.50 0.8390332 0.6749413
## 0.75 0.8391053 0.6752263
## 1.00 0.8391053 0.6753687
## 1.25 0.8391053 0.6753687
## 1.50 0.8406205 0.6783406
## 1.75 0.8422078 0.6815951
## 2.00 0.8406205 0.6785395
## 5.00 0.8453102 0.6879521
##
## Accuracy was used to select the optimal model using the largest value.
## The final value used for the model was C = 5.
```

```
# accuracy plot of tuned model
plot(svm.mod.grid)
```



```
# prediction using tuned model
svm.pred.grid <- predict(svm.mod.grid, newdata = testing)
svm.pred.grid
```

```
## [1] 1 1 1 1 1 1 1 0 1 1 1 1 1 1 1 0 0 0 1 0 1 1 1 1 1 1 1 1 1 0 0 1 1 1 1 1
## [39] 1 1 1 1 1 1 1 0 0 1 1 1 1 0 1 0 1 1 1 0 0 0 0 0 0 0 0 0 0 0 1 0 0 0 0 0 0 0
## [77] 0 0 0 1 1 1 0 0 1 0 1 0 0 0
## Levels: 0 1
```

```
# accuracy of the tuned model
confusionMatrix(table(svm.pred.grid, testing$target))
```

```
## Confusion Matrix and Statistics
##
##
## svm.pred.grid  0  1
##              0 28 10
##              1  9 43
##
##              Accuracy : 0.7889
##              95% CI : (0.6901, 0.8679)
##      No Information Rate : 0.5889
##      P-Value [Acc > NIR] : 4.918e-05
##
##              Kappa : 0.5658
##
##  Mcnemar's Test P-Value : 1
##
##              Sensitivity : 0.7568
##              Specificity : 0.8113
##      Pos Pred Value : 0.7368
##      Neg Pred Value : 0.8269
##      Prevalence : 0.4111
##      Detection Rate : 0.3111
##      Detection Prevalence : 0.4222
##      Balanced Accuracy : 0.7840
##
##      'Positive' Class : 0
##
```

Logistic Regression

```
set.seed(2441139)
# Organize data to get training and test data

X <- as.matrix(heart, c("age", "sex", "cp", "trestbps", "chol", "fbs",
                        "restecg", "thalach", "exang", "oldpeak", "slope",
                        "ca", "thal"))
```

```
## Warning in if (rownames.force %in% FALSE) NULL else if (rownames.force %in% :
## the condition has length > 1 and only the first element will be used
```

```
y <- heart$target

n <- nrow(X)
train_rows <- sample(1:n, n * 0.7)

X.train <- X[train_rows,]
X.test <- X[-train_rows,]
```



```
y.train <- y[train_rows]
y.test <- y[-train_rows]
```

```
dim(X.train)
```

```
## [1] 212 14
```

```
dim(X.test)
```

```
## [1] 91 14
```

```
grid <- 10^seq(10, -2, length = 100)
```


```
# ----- #
#                               Lasso                               #
# ----- #
```

```
# lasso model
```

```
lasso.mod <- glmnet(X.train, as.factor(y.train), alpha = 1, lambda = grid,
                    family = "binomial")
```

```
plot(lasso.mod, xvar = "lambda", label = T)
```

```
## Warning in plotCoef(x$beta, lambda = x$lambda, df = x$df, dev = x$dev.ratio, : 1
## or less nonzero coefficients; glmnet plot is not meaningful
```



```
# cross-validation for lambda
cv.out <- cv.glmnet(X.train, as.factor(y.train), family = "binomial", alpha = 1,
                    type.measure = "class")
bestlam <- cv.out$lambda.min
bestlam
```

```
## [1] 0.4144432
```

```
# coefficients of the best model
best.lasso.mod <- glmnet(X.train, as.factor(y.train), alpha = 1, lambda = bestlam,
                        family = "binomial")
coef(best.lasso.mod)
```

```
## 15 x 1 sparse Matrix of class "dgCMatrix"
##                               s0
## (Intercept) -0.2468234
## i.age       .
## sex         .
## cp          .
## trestbps    .
## chol        .
## fbs         .
## restecg     .
## thalach     .
## exang       .
## oldpeak     .
## slope       .
## ca          .
## thal        .
## target      0.6873644
```


```
# test error
lasso.pred <- predict(best.lasso.mod, newx = X.test, s = bestlam)
lasso.mse <- mean((lasso.pred - y.test)^2)
lasso.mse
```

```
## [1] 0.2077333
```

```
# non-zero coefficients
lasso.coef <- predict(best.lasso.mod, type = "coefficients", s = bestlam)
lasso.coef <- lasso.coef[which(lasso.coef != 0)]
lasso.coef
```

```
## [1] -0.2468234  0.6873644
```

```
# ----- #
#                               Elastic Net                               #
# ----- #
# elastic net model
en.mod <- glmnet(X.train, as.factor(y.train), alpha = 0.5, lambda = grid,
                family = "binomial")
plot(en.mod, xvar = "lambda", label = T)
```



```
# cross-validation for lambda (with a fixed alpha)
cv.out <- cv.glmnet(X.train, y.train, alpha = 0.5)
bestlam <- cv.out$lambda.min

# coefficients of the best model
best.en.mod <- glmnet(X.train, as.factor(y.train), alpha = 0.5, lambda = bestlam,
                      family = "binomial")
coef(best.en.mod)
```

```
## 15 x 1 sparse Matrix of class "dgCMatrix"
##                               s0
## (Intercept) -3.323601000
## i..age      .
## sex         -0.262775524
## cp          0.141658244
## trestbps    .
## chol        .
## fbs         .
## restecg     .
## thalach     0.006374926
## exang       -0.286353082
## oldpeak     -0.104291670
## slope       0.033116852
## ca         -0.122549553
## thal        -0.172448531
## target      6.221297686
```

```
# test error
en.pred <- predict(best.en.mod, s = bestlam, newx = X.test)
en.mse <- mean((en.pred - y.test)^2)
en.mse
```

```
## [1] 8.411157
```

```
# non-zero coefficients
en.coef <- predict(best.en.mod, type = "coefficients", s = bestlam)
en.coef <- en.coef[which(en.coef != 0)]
en.coef
```

```
## [1] -3.323601000 -0.262775524 0.141658244 0.006374926 -0.286353082
## [6] -0.104291670 0.033116852 -0.122549553 -0.172448531 6.221297686
```