

# CS5525 Final Project Code Submission

Jennifer Appiah-Kubi, Rebecca DeSipio, Ajinkya Fotedar

12/11/2021

## Contents

<b>Classification Methods</b>	<b>1</b>
Decision Trees . . . . .	1
KNN . . . . .	8
SVM . . . . .	9
Logistic Regression . . . . .	13

## Classification Methods

First set the directory (path which contains the `heart.csv` data), and import any needed libraries

```
#setwd(" ") # uncomment to set working directory via code
library(tree)
library(randomForest) # bootstrap/bagging & random forest
library(class)        # KNN
library(caret)         # SVM
library(ggplot2)
library(cowplot)
```

## Decision Trees

```
# Read in and organize data
## -- Read data
heart <- read.csv("heart.csv")
Target <- as.factor(heart$target) # target heart rate

## -- Split into training and test sets
train <- sample(1:nrow(heart), 0.75*nrow(heart))
heart.test <- heart[-train, ]
Target.test <- Target[-train]

# ----- #
#                               Fit a Classification Tree                               #
# ----- #


# Fit a classification tree to the training data
set.seed(2441139)
tree.heart <- tree(Target~. -target, heart, subset=train)
summary(tree.heart)

##
```

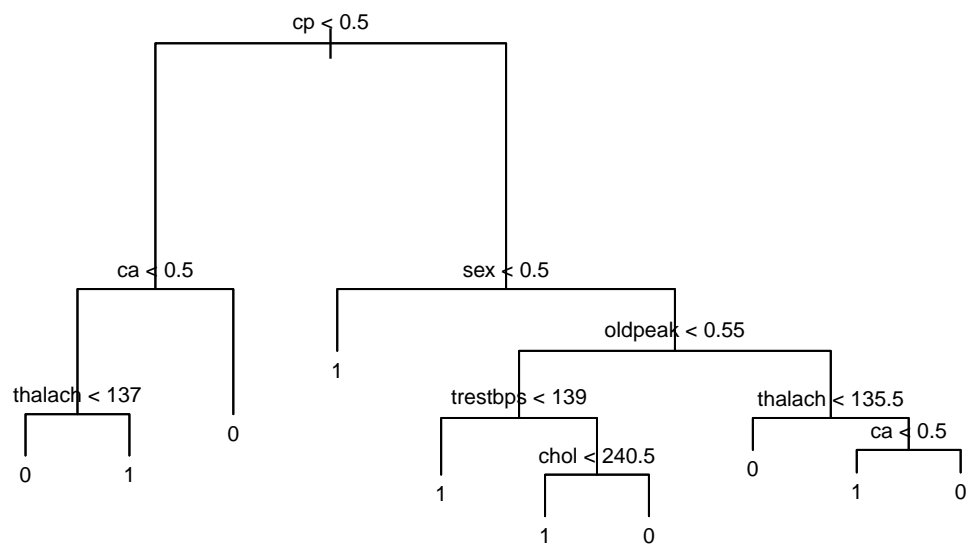
```

graph TD
    Root["cp < 0.5"]
    Root --> Left["ca < 0.5"]
    Root --> Right["sex < 0.5"]
    
    Left --> Left_L["thalach < 137"]
    Left --> Left_R["oldpeak < 0.4"]
    
    Left_L --> Left_L_L["oldpeak < 1.5"]
    Left_L --> Left_L_R["exang < 0.5"]
    
    Left_L_L --> Left_L_L_L["0"]
    Left_L_L --> Left_L_L_R["0"]
    
    Left_L_R --> Left_L_R_L["1"]
    Left_L_R --> Left_L_R_R["1"]
    
    Left_R --> Left_R_L["thalach < 160"]
    Left_R --> Left_R_R["chol < 301"]
    
    Left_R_L --> Left_R_L_L["0"]
    Left_R_L --> Left_R_L_R["0"]
    
    Left_R_R --> Left_R_R_L["0"]
    Left_R_R --> Left_R_R_R["0"]
    
    Right --> Right_L["ca < 0.5"]
    Right --> Right_R["oldpeak < 0.55"]
    
    Right_L --> Right_L_L["age < 62.5"]
    
    Right_L_L --> Right_L_L_L["1"]
    Right_L_L --> Right_L_L_R["1"]
    
    Right_L_R --> Right_L_R_L["1"]
    
    Right_R --> Right_R_L["trestbps < 139"]
    Right_R --> Right_R_R["thalach < 135.5"]
    
    Right_R_L --> Right_R_L_L["chol < 240.5"]
    
    Right_R_L_L --> Right_R_L_L_L["1"]
    Right_R_L_L --> Right_R_L_L_R["1"]
    
    Right_R_L_R --> Right_R_L_R_L["0"]
    
    Right_R_R --> Right_R_R_L["slope < 1.5"]
    
    Right_R_R_L --> Right_R_R_L_L["0"]
    Right_R_R_L --> Right_R_R_L_R["1"]
    
    Right_R_R_L_R --> Right_R_R_R["1"]
    Right_R_R_R --> Right_R_R_R_L["1"]
    Right_R_R_R --> Right_R_R_R_R["1"]
  
```

```
## $size
## [1] 18 10 8 4 2 1
##
## $dev
## [1] 52 52 53 54 63 98
##
## $k
## [1] -Inf 0.00 2.00 2.75 7.00 43.00
##
## $method
## [1] "misclass"
##
## attr("class")
## [1] "prune" "tree.sequence"
plot(cv.heart$size, cv.heart$dev, type='b')
```



```
prune.heart <- prune.misclass(tree.heart, best=10)
plot(prune.heart)
text(prune.heart, pretty=1, cex=0.65)
```



```
# Predict using test set and pruned tree. Compare.
tree.pred <- predict(tree.heart, heart.test, type='class') # test tree
prune.pred <- predict(prune.heart, heart.test, type='class') # pruned tree
```

```
table(prune.pred, Target.test)
```

```
##           Target.test
## prune.pred  0  1
##           0 25  9
##           1 17 25
```

```
table(tree.pred, Target.test)
```

```
##           Target.test
## tree.pred 0  1
##           0 25  9
##           1 17 25
```

```
# -----#
#                               Bagging                               #
# -----#
```

```
set.seed(2441139)
```

```
# Perform bagging
```

```
bag.heart <- randomForest(as.factor(as.character(heart$target))~., data=heart,
                          subset=train, mtry=ncol(heart)-1,
                          importance=TRUE)
```

```
bag.heart
```

```
##
```

```
## Call:
```

```
## randomForest(formula = as.factor(as.character(heart$target)) ~ ., data = heart, mtry = ncol(heart))
```

```
##           Type of random forest: classification
```

```
##           Number of trees: 500
```

```
## No. of variables tried at each split: 13
```

```
##
```

```
##           OOB estimate of  error rate: 18.5%
```

```
## Confusion matrix:
```

```
##           0  1 class.error
```

```
## 0 71  25  0.2604167
```

```
## 1 17 114  0.1297710
```

```
# Predict on bagged tree
```

```
bag.pred <- predict(bag.heart, heart.test, type='class')
```

```
table(bag.pred, Target.test)
```

```
##           Target.test
```


```
## bag.pred 0  1
```

```
##           0 30  5
```

```
##           1 12 29
```

```
varImpPlot(bag.heart)
```

## bag.heart



```
# -----#
#                               Random Forest                               #
# -----#
```

```
set.seed(2441139)
```

```
# Perform Random Forest
```

```
rf.heart <- randomForest(as.factor(as.character(heart$target))~., data=heart,
                          subset=train, mtry=sqrt(ncol(heart)-1),
                          ntree=25, importance=TRUE)
```

```
rf.heart
```

```
##
```

```
## Call:
```

```
## randomForest(formula = as.factor(as.character(heart$target)) ~ ., data = heart, mtry = sqrt(ncol(heart)-1),
```

```
##               Type of random forest: classification
```

```
##               Number of trees: 25
```

```
## No. of variables tried at each split: 4
```

```
##
```

```
## OOB estimate of error rate: 20.26%
```

```
## Confusion matrix:
```

```
##    0  1 class.error
```

```
## 0 66 30  0.3125000
```

```
## 1 16 115  0.1221374
```

```
# Predict on the forest
```

```
rf.pred <- predict(rf.heart, heart.test, type='class')
```

```
table(rf.pred, Target.test)
```


```
## Target.test
```

```
## rf.pred 0 1
```


```
##      0 29  3
##      1 13 31
```

```
varImpPlot(rf.heart)
```

rf.heart



```
# ----- #
# ----- #
# Determine Best Model (Random Forest) #
# ----- #
# ----- #
# Investigate how mtry affect the accuracy
Acc <- rep(0,ncol(heart)-2)
for (m in 1:(ncol(heart)-2)){
  set.seed(2441139)
  rf.heart <- randomForest(as.factor(as.character(heart$target))~., data=heart,
                           subset=train, mtry=m,
                           ntree=25)
  rf.pred <- predict(rf.heart, heart.test, type='class')
  t <- table(rf.pred, Target.test)
  acc <- sum(diag(t))/sum(t)
  Acc[m] <- acc
}
mbest <- which(Acc==max(Acc))
plot(1:(ncol(heart)-2), Acc, xlab='mtry', ylab='Accuracy of random forest') # include plot in final sub
```



```
# Now use the best value of m for the random forest
```

```
set.seed(2441139)
```

```
rf.heart <- randomForest(as.factor(as.character(heart$target))~., data=heart,  
                          subset=train, mtry=mbest,  
                          ntree=25, importance=TRUE)
```

```
rf.heart
```

```
##
```

```
## Call:
```

```
## randomForest(formula = as.factor(as.character(heart$target)) ~ ., data = heart, mtry = mbest, ntree = 25, importance = TRUE)
```

```
##           Type of random forest: classification
```

```
##           Number of trees: 25
```

```
## No. of variables tried at each split: 2
```

```
##
```

```
##           OOB estimate of  error rate: 17.62%
```

```
## Confusion matrix:
```

```
##      0   1 class.error
```

```
## 0  77  19  0.1979167
```

```
## 1  21 110  0.1603053
```

```
# Predict on the forest
```

```
rf.pred <- predict(rf.heart, heart.test, type='class')
```

```
table(rf.pred, Target.test)
```

```
##           Target.test
```


```
## rf.pred  0  1
```

```
##           0 31  7
```

```
##           1 11 27
```

```
varImpPlot(rf.heart)
```

## rf.heart



## KNN

```
# K-Nearest Neighbor
set.seed(2441139)

# Read data
heart <- read.csv("heart.csv")
Target <- as.factor(heart$target)

# Split into training and test sets
train <- sample(1:nrow(heart), 0.75*nrow(heart))
heart.test <- heart[-train, ]
Target.test <- Target[-train]

# Normalizing function
normalizer <-function(x){(x -min(x))/(max(x)-min(x))}

# Use normalized set to do cross-validation
trControl <- trainControl(method = "cv",
                           number = 5)
fit <- train(as.factor(target) ~ .,
             method = "knn",
             tuneGrid = expand.grid(k = 1:15),
             trControl = trControl,
             metric = "Accuracy",
             data = as.data.frame(lapply(heart, normalizer)))

# For best value of k perform training
```



```

heart.norm <- as.data.frame(lapply(heart[,1:13], normalizer))

# Get training and test set
heart.norm.train <- heart.norm[train,]
heart.norm.test <- heart.norm[-train,]

# Get cl argument
cl <- as.factor(heart$target[train])

# Train and get confusion matrix
knn.heart <- knn(heart.norm.train, heart.norm.test, cl, k = 12, prob=TRUE)
cm <- table(knn.heart, Target.test)
(cm)

##           Target.test
## knn.heart  0  1
##           0 24 11
##           1  5 36

acc <- sum(diag(cm))/sum(cm)

# Train and get confusion matrix for un-normalized data
knn.heart <- knn(heart[train,1:13],heart[-train,1:13], cl, k = 5, prob=TRUE)
cm <- table(knn.heart, Target.test)
(cm)

##           Target.test
## knn.heart  0  1
##           0 19 14
##           1 10 33

acc <- sum(diag(cm))/sum(cm)

```

## SVM

```

set.seed(2441139)

# splitting data into test and train
intrain <- createDataPartition(y = heart$target, p = 0.7, list = F)

training <- heart[intrain,]
testing <- heart[-intrain,]
training[["target"]] <- as.factor(training[["target"]])

dim(training)

## [1] 213 14

dim(testing)

## [1] 90 14

# set up repeated k-fold cross-validation
trctrl <- trainControl(method = "repeatedcv", number = 10, repeats = 3)

# model training with svm

```

```

svm.mod <- train(target ~ ., data = training, method = "svmLinear",
                 trControl = trctrl,
                 preProcess = c("center", "scale"),
                 tuneLength = 10)

svm.mod

## Support Vector Machines with Linear Kernel
##
## 213 samples
## 13 predictor
## 2 classes: '0', '1'
##
## Pre-processing: centered (13), scaled (13)
## Resampling: Cross-Validated (10 fold, repeated 3 times)
## Summary of sample sizes: 192, 191, 192, 192, 191, 191, ...
## Resampling results:
##
## Accuracy Kappa
## 0.8266955 0.6511675
##
## Tuning parameter 'C' was held constant at a value of 1

# prediction using the above model
svm.pred <- predict(svm.mod, newdata = testing)
svm.pred

## [1] 1 1 1 1 1 1 1 0 1 1 1 1 1 1 1 1 0 0 0 1 0 1 1 1 1 1 1 0 1 1 0 0 1 1 1 1 1
## [39] 1 1 1 1 1 1 1 0 0 1 1 1 1 0 1 0 1 1 1 0 0 0 0 0 0 0 0 0 0 0 1 0 0 0 0 0 0 0
## [77] 0 0 0 1 1 1 0 0 1 0 1 0 0 0
## Levels: 0 1

# accuracy of the trained model
confusionMatrix(table(svm.pred, testing$target))

## Confusion Matrix and Statistics
##
##
## svm.pred 0 1
##      0 28 11
##      1 9 42
##
##              Accuracy : 0.7778
##              95% CI : (0.6779, 0.8587)
##      No Information Rate : 0.5889
##      P-Value [Acc > NIR] : 0.0001266
##
##              Kappa : 0.5448
##
## Mcnemar's Test P-Value : 0.8230633
##
##              Sensitivity : 0.7568
##              Specificity : 0.7925
##      Pos Pred Value : 0.7179
##      Neg Pred Value : 0.8235
##              Prevalence : 0.4111

```

```

##          Detection Rate : 0.3111
##    Detection Prevalence : 0.4333
##          Balanced Accuracy : 0.7746
##
##          'Positive' Class : 0
##

# costs for further tuning
grid <- expand.grid(C = seq(0, 2, length = 20))


svm.mod.tuned <- train(target ~ ., data = training, method = "svmLinear",
                      trControl = trctrl,
                      preProcess = c("center", "scale"),
                      tuneGrid = grid,
                      tuneLength = 10)

svm.mod.tuned

## Support Vector Machines with Linear Kernel
##
## 213 samples
## 13 predictor
## 2 classes: '0', '1'
##
## Pre-processing: centered (13), scaled (13)
## Resampling: Cross-Validated (10 fold, repeated 3 times)
## Summary of sample sizes: 192, 191, 192, 192, 192, 192, ...
## Resampling results across tuning parameters:
##
##  C           Accuracy   Kappa
##  0.0000000      NaN      NaN
##  0.1052632  0.8421356  0.6805929
##  0.2105263  0.8452381  0.6870639
##  0.3157895  0.8421356  0.6810364
##  0.4210526  0.8437229  0.6842502
##  0.5263158  0.8422078  0.6814217
##  0.6315789  0.8406205  0.6783419
##  0.7368421  0.8391053  0.6752263
##  0.8421053  0.8375180  0.6721433
##  0.9473684  0.8391053  0.6753687
##  1.0526316  0.8391053  0.6753687
##  1.1578947  0.8391053  0.6753687
##  1.2631579  0.8375180  0.6721433
##  1.3684211  0.8375180  0.6721433
##  1.4736842  0.8390332  0.6751736
##  1.5789474  0.8422078  0.6815951
##  1.6842105  0.8422078  0.6815951
##  1.7894737  0.8406205  0.6785395
##  1.8947368  0.8406205  0.6785395
##  2.0000000  0.8406205  0.6785395
##
## Accuracy was used to select the optimal model using the largest value.
## The final value used for the model was C = 0.2105263.

# accuracy plot with varying costs
plot(svm.mod.tuned)

```



```
# best cost value that maximizes model accuracy
```

```
svm.mod.tuned$bestTune
```

```
##          C
```

```
## 3 0.2105263
```

```
# prediction using tuned model
```

```
svm.pred.tuned <- predict(svm.mod.tuned, newdata = testing)
```

```
svm.pred.tuned
```

```
## [1] 1 1 1 1 1 1 0 1 1 1 1 1 1 1 1 1 1 1 0 1 1 1 1 1 1 1 1 1 1 1 0 1 1 1 1 1
```

```
## [39] 1 1 1 1 1 1 1 1 0 1 1 1 1 0 1 0 1 1 1 0 0 0 0 0 0 0 0 0 0 0 1 0 0 0 0 0 0 0
```

```
## [77] 0 0 0 1 1 1 0 0 1 0 1 0 0 0
```

```
## Levels: 0 1
```

```
# accuracy of the tuned model
```

```
confusionMatrix(table(svm.pred.tuned, testing$target))
```

```
## Confusion Matrix and Statistics
```

```
##
```

```
##
```

```
## svm.pred.tuned  0  1
```

```
##              0 28  5
```

```
##              1  9 48
```

```
##
```

```
##              Accuracy : 0.8444
```

```
##              95% CI : (0.7528, 0.9123)
```

```
##              No Information Rate : 0.5889
```

```
##              P-Value [Acc > NIR] : 1.588e-07
```

```
##
```

```
##              Kappa : 0.6734
```

```
##
```

```
## McNemar's Test P-Value : 0.4227
##
##          Sensitivity : 0.7568
##          Specificity : 0.9057
##          Pos Pred Value : 0.8485
##          Neg Pred Value : 0.8421
##          Prevalence : 0.4111
##          Detection Rate : 0.3111
##          Detection Prevalence : 0.3667
##          Balanced Accuracy : 0.8312
##
##          'Positive' Class : 0
##
```

## Logistic Regression

```
set.seed(2441139)
```

```
# pre-processing of data
str(heart)
```

```
## 'data.frame': 303 obs. of 14 variables:
## $ age : int 63 37 41 56 57 57 56 44 52 57 ...
## $ sex : int 1 1 0 1 0 1 0 1 1 1 ...
## $ cp : int 3 2 1 1 0 0 1 1 2 2 ...
## $ trestbps: int 145 130 130 120 120 140 140 120 172 150 ...
## $ chol : int 233 250 204 236 354 192 294 263 199 168 ...
## $ fbs : int 1 0 0 0 0 0 0 0 1 0 ...
## $ restecg : int 0 1 0 1 1 1 0 1 1 1 ...
## $ thalach : int 150 187 172 178 163 148 153 173 162 174 ...
## $ exang : int 0 0 0 0 1 0 0 0 0 0 ...
## $ oldpeak : num 2.3 3.5 1.4 0.8 0.6 0.4 1.3 0 0.5 1.6 ...
## $ slope : int 0 0 2 2 2 1 1 2 2 2 ...
## $ ca : int 0 0 0 0 0 0 0 0 0 0 ...
## $ thal : int 1 2 2 2 2 1 2 3 3 2 ...
## $ target : int 1 1 1 1 1 1 1 1 1 1 ...
```

```
heart$sex <- ifelse(test = heart$sex == 0, yes = "F", no = "M")
heart$sex <- as.factor(heart$sex)
heart$cp <- as.factor(heart$cp)
heart$fbs <- as.factor(heart$fbs)
heart$restecg <- as.factor(heart$restecg)
heart$exang <- as.factor(heart$exang)
heart$slope <- as.factor(heart$slope)
heart$ca <- as.factor(heart$ca)
heart$thal <- as.factor(heart$thal)
heart$age <- as.numeric(heart$age)
heart$trestbps <- as.numeric(heart$trestbps)
heart$chol <- as.numeric(heart$chol)
heart$thalach <- as.numeric(heart$thalach)
heart$target <- ifelse(test = heart$target == 0, yes = "Healthy", no = "Unhealthy")
heart$target <- as.factor(heart$target)

str(heart)
```

```
## 'data.frame': 303 obs. of 14 variables:
## $ age : num 63 37 41 56 57 57 56 44 52 57 ...
## $ sex : Factor w/ 2 levels "F","M": 2 2 1 2 1 2 1 2 2 2 ...
## $ cp : Factor w/ 4 levels "0","1","2","3": 4 3 2 2 1 1 2 2 3 3 ...
## $ trestbps: num 145 130 130 120 120 140 140 120 172 150 ...
## $ chol : num 233 250 204 236 354 192 294 263 199 168 ...
## $ fbs : Factor w/ 2 levels "0","1": 2 1 1 1 1 1 1 1 2 1 ...
## $ restecg : Factor w/ 3 levels "0","1","2": 1 2 1 2 2 2 1 2 2 2 ...
## $ thalach : num 150 187 172 178 163 148 153 173 162 174 ...
## $ exang : Factor w/ 2 levels "0","1": 1 1 1 1 2 1 1 1 1 1 ...
## $ oldpeak : num 2.3 3.5 1.4 0.8 0.6 0.4 1.3 0 0.5 1.6 ...
## $ slope : Factor w/ 3 levels "0","1","2": 1 1 3 3 3 2 2 3 3 3 ...
## $ ca : Factor w/ 5 levels "0","1","2","3",...: 1 1 1 1 1 1 1 1 1 1 ...
## $ thal : Factor w/ 4 levels "0","1","2","3": 2 3 3 3 3 2 3 4 4 3 ...
## $ target : Factor w/ 2 levels "Healthy","Unhealthy": 2 2 2 2 2 2 2 2 2 2 ...
```

```
# getting the number of samples by gender
xtabs(~ target + sex, data = heart)
```

```
##           sex
## target      F  M
##   Healthy   24 114
##   Unhealthy  72  93
```

```
# getting the number of samples by chest pain
xtabs(~ target + cp, data = heart)
```

```
##           cp
## target      0  1  2  3
##   Healthy  104  9 18  7
##   Unhealthy  39 41 69 16
```

```
# getting the number of samples by resting ECG
xtabs(~ target + restecg, data = heart)
```

```
##           restecg
## target      0  1  2
##   Healthy   79 56  3
##   Unhealthy  68 96  1
```

```
# simple logistic model
logistic <- glm(target ~ sex, data = heart, family = "binomial")
summary(logistic)
```

```
##
## Call:
## glm(formula = target ~ sex, family = "binomial", data = heart)
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -1.6651  -1.0923   0.7585   1.2650   1.2650
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)    1.0986     0.2357   4.661 3.15e-06 ***
## sexM           -1.3022     0.2740  -4.752 2.01e-06 ***
## ---
```

```

## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##      Null deviance: 417.64  on 302  degrees of freedom
## Residual deviance: 392.80  on 301  degrees of freedom
## AIC: 396.8
##
## Number of Fisher Scoring iterations: 4
R_sq_1 <- 1 - logistic$deviance / logistic$null.deviance
R_sq_1

## [1] 0.05947945
BIC_1 <- logistic$deviance + 2 * log(dim(heart)[1])
BIC_1

## [1] 404.2246
# complex logistic model
logistic <- glm(target ~ ., data = heart, family = "binomial")
summary(logistic)

##
## Call:
## glm(formula = target ~ ., family = "binomial", data = heart)
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -2.9459  -0.2738   0.1012   0.4515   3.1248
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)  0.179045   3.705420   0.048  0.961461
## age          0.027819   0.025428   1.094  0.273938
## sexM        -1.862297   0.570844  -3.262  0.001105 **
## cp1          0.864708   0.578000   1.496  0.134645
## cp2          2.003186   0.529356   3.784  0.000154 ***
## cp3          2.417107   0.719242   3.361  0.000778 ***
## trestbps    -0.026162   0.011943  -2.191  0.028481 *
## chol        -0.004291   0.004245  -1.011  0.312053
## fbs1         0.445666   0.587977   0.758  0.448472
## restecg1     0.460582   0.399615   1.153  0.249089
## restecg2    -0.714204   2.768873  -0.258  0.796453
## thalach      0.020055   0.011859   1.691  0.090820 .
## exang1      -0.779111   0.451839  -1.724  0.084652 .
## oldpeak     -0.397174   0.242346  -1.639  0.101239
## slope1      -0.775084   0.880495  -0.880  0.378707
## slope2       0.689965   0.947657   0.728  0.466568
## ca1         -2.342301   0.527416  -4.441  8.95e-06 ***
## ca2         -3.483178   0.811640  -4.292  1.77e-05 ***
## ca3         -2.247144   0.937629  -2.397  0.016547 *
## ca4          1.267961   1.720014   0.737  0.461013
## thal1        2.637558   2.684285   0.983  0.325808
## thal2        2.367747   2.596159   0.912  0.361759

```

```

## thal3      0.915115   2.600380   0.352 0.724901
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##      Null deviance: 417.64  on 302  degrees of freedom
## Residual deviance: 179.63  on 280  degrees of freedom
## AIC: 225.63
##
## Number of Fisher Scoring iterations: 6
R_sq_2 <- 1 - logistic$deviance / logistic$null.deviance
R_sq_2

## [1] 0.569889
BIC_2 <- logistic$deviance + 14 * log(dim(heart)[1])
BIC_2


## [1] 259.623
# why age isn't of statistical significance
median(heart$age)

## [1] 55
# plotting the probability of getting a heart disease
predict.hd <- data.frame(prob.of.hd = logistic$fitted.values, hd = heart$target)
predict.hd <- predict.hd[order(predict.hd$prob.of.hd, decreasing = FALSE), ]
predict.hd$rank <- 1:nrow(predict.hd)

ggplot(data = predict.hd, aes(x = rank, y = prob.of.hd)) +
  geom_point(aes(color = hd), alpha = 1, shape = 4, stroke = 2) +
  xlab("Index") +
  ylab("Predicted probability of getting a heart disease")

```





```
# predicting model accuracy
log.mod <- glm(target ~ ., data = training, family = "binomial")
log.pred <- predict(log.mod, newdata = testing, type = "response")
log.pred <- ifelse(log.pred > 0.5, 1, 0)
confusionMatrix(table(log.pred, testing$target))
```

```
## Confusion Matrix and Statistics
##
##
## log.pred  0  1
##          0 28 11
##          1  9 42
##
##               Accuracy : 0.7778
##               95% CI : (0.6779, 0.8587)
##      No Information Rate : 0.5889
##      P-Value [Acc > NIR] : 0.0001266
##
##               Kappa : 0.5448
##
##  Mcnemar's Test P-Value : 0.8230633
##
##               Sensitivity : 0.7568
##               Specificity : 0.7925
##      Pos Pred Value : 0.7179
##      Neg Pred Value : 0.8235
##      Prevalence : 0.4111
```

```
##          Detection Rate : 0.3111
## Detection Prevalence : 0.4333
##    Balanced Accuracy : 0.7746
##
##    'Positive' Class : 0
##
```